

INTRODUCTION TO OPENACC

Profiling, Parallelizing, and Optimizing with OpenACC



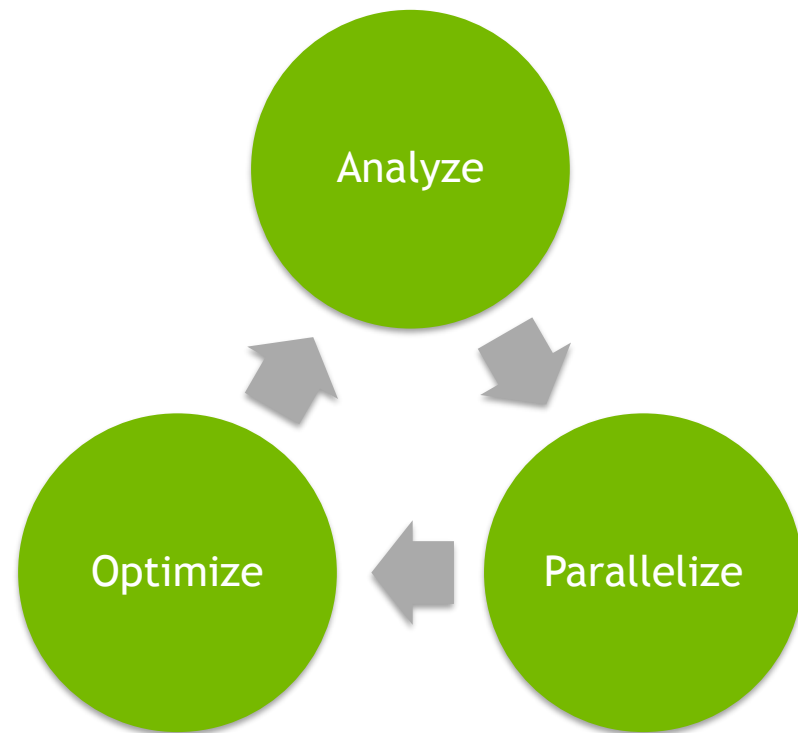
Objectives

Understand what OpenACC is and why to use it

Learn how to obtain an application profile using PGProf

Learn how to add OpenACC directives to existing loops and build with OpenACC using PGI

Perform simple data and loop optimizations to improve performance





Why OpenACC?

OpenACC

Simple | Powerful | Portable

Fueling the Next Wave of
Scientific Discoveries in HPC

```
main()
{
    <serial code>
    #pragma acc kernels
    //automatically runs on GPU
    {
        <parallel code>
    }
}
```

University of Illinois
PowerGrid- MRI Reconstruction



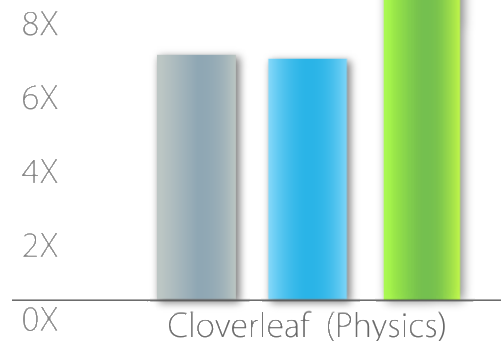
70x Speed-Up
2 Days of Effort

RIKEN Japan
NICAM- Climate Modeling



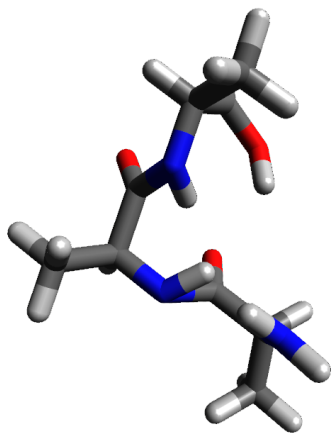
7-8x Speed-Up
5% of Code Modified

■ CPU: OpenMP
■ CPU: OpenACC
■ GPU: OpenACC



LS-DALTON

Large-scale application for calculating high-accuracy molecular energies



“OpenACC makes GPU computing approachable for domain scientists. Initial OpenACC implementation required only minor effort, and more importantly, *no modifications* of our existing CPU implementation.”

Janus Juul Eriksen, PhD Fellow
qLEAP Center for Theoretical Chemistry, Aarhus University



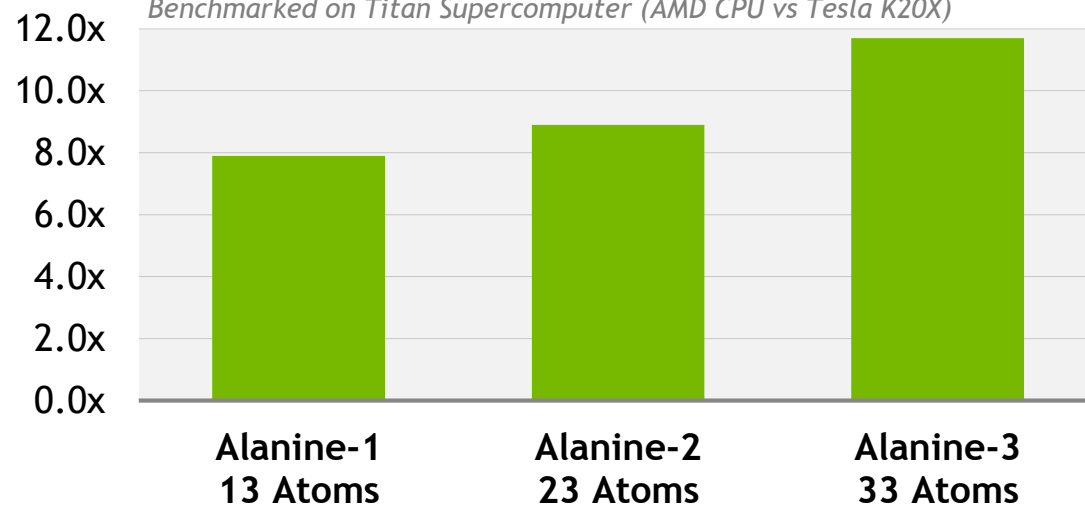
Minimal Effort

Lines of Code Modified	# of Weeks Required	# of Codes to Maintain
<100 Lines	1 Week	1 Source

Big Performance

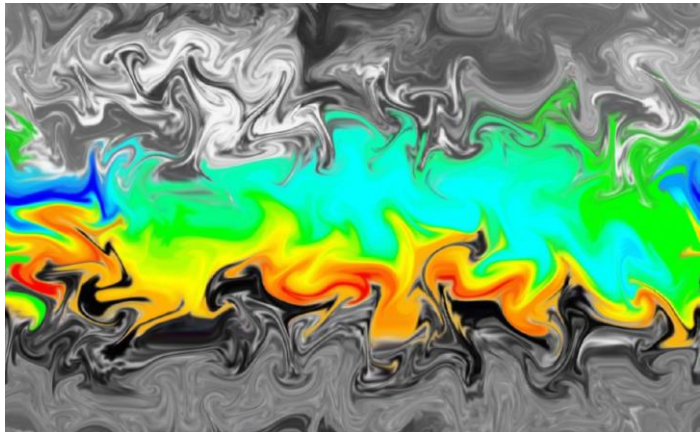
LS-DALTON CCSD(T) Module

Benchmarked on Titan Supercomputer (AMD CPU vs Tesla K20X)



OpenACC Performance Portability: CloverLeaf

Hydrodynamics Application

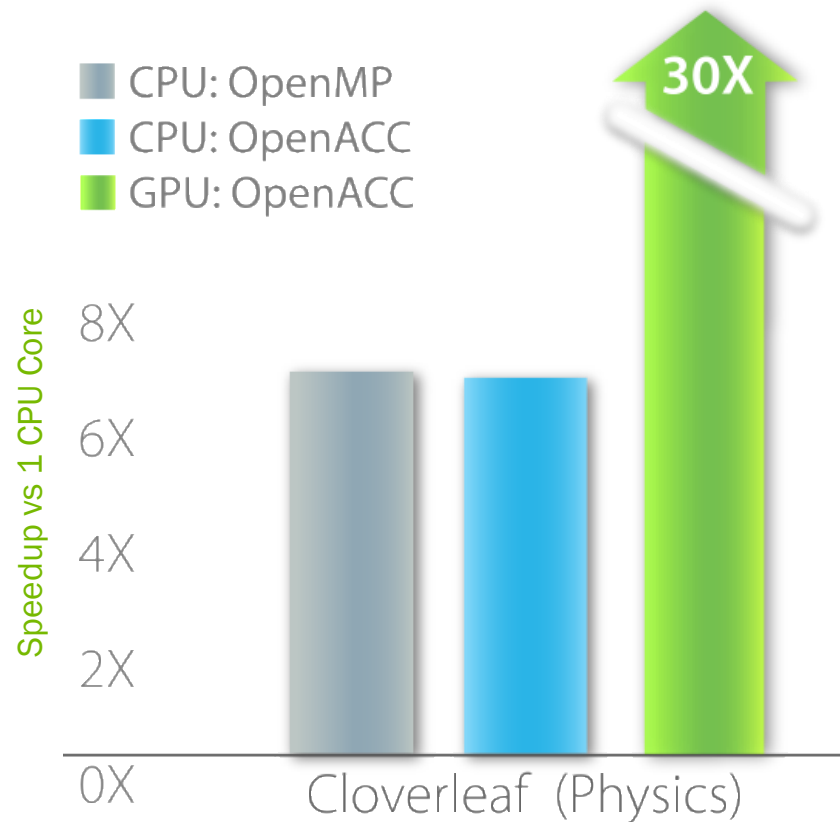


“We were extremely impressed that we can run OpenACC on a CPU with *no code change* and get *equivalent performance* to our OpenMP/MPI implementation.”

Wayne Gaudin and Oliver Perks
Atomic Weapons Establishment, UK

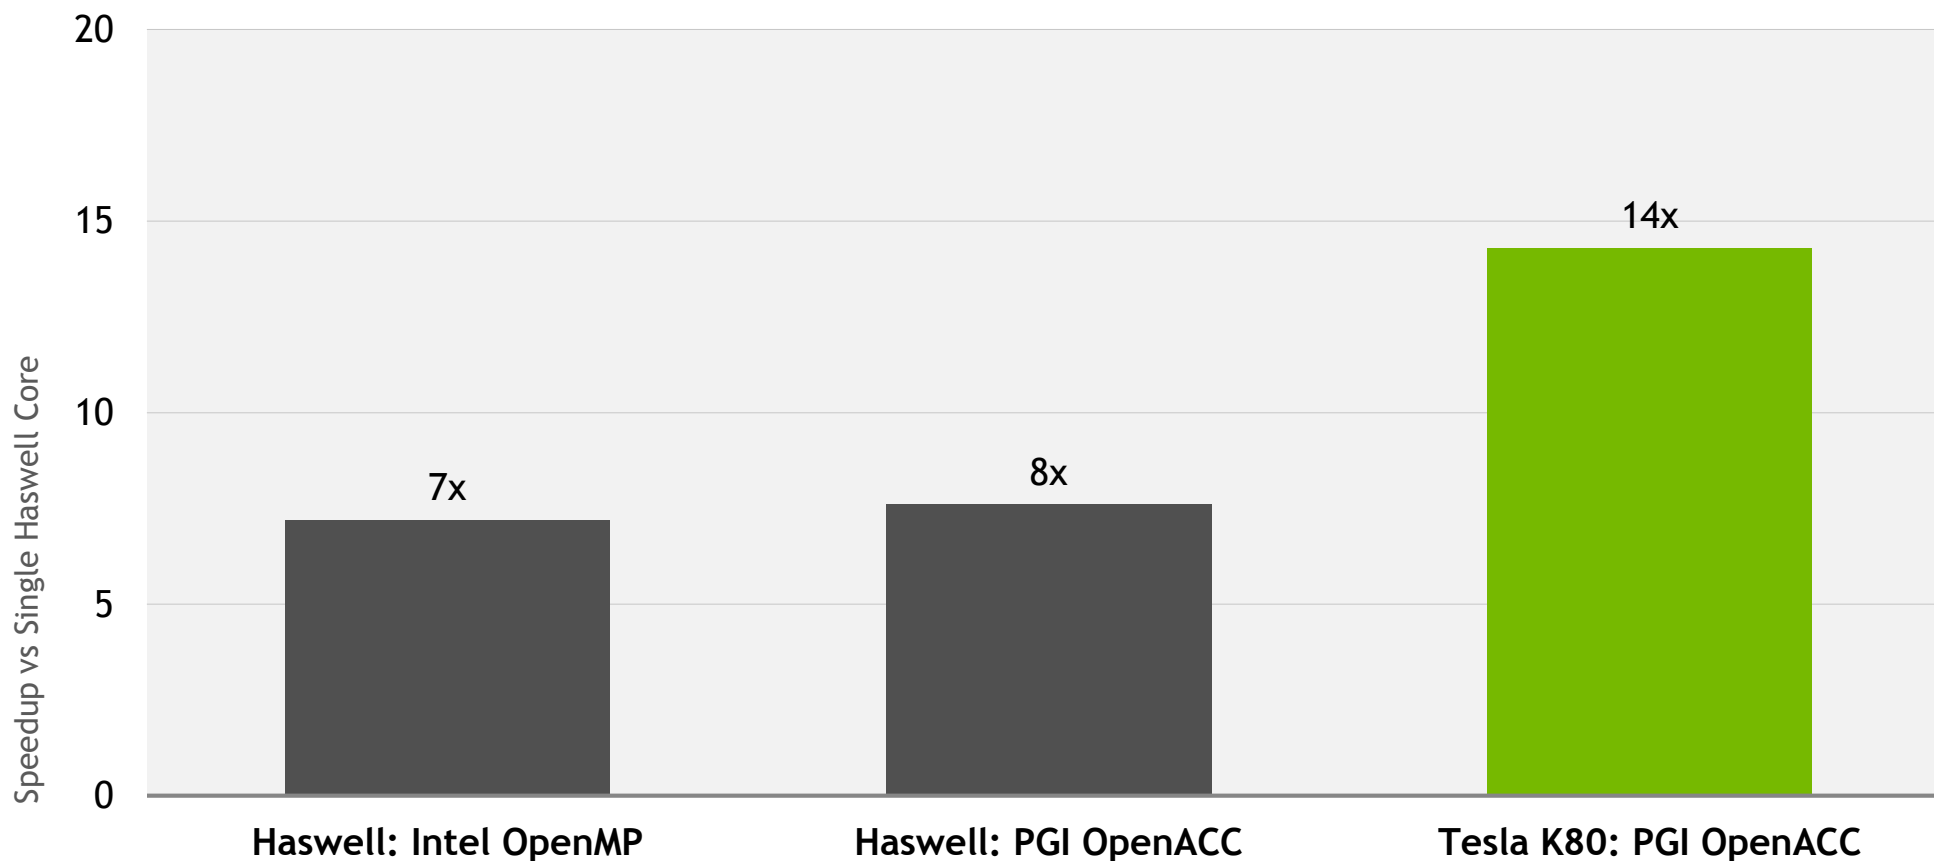


OpenACC Performance Portability



Benchmarked Intel(R) Xeon(R) CPU E5-2690 v2 @ 3.00GHz, Accelerator: Tesla K80

CloverLeaf on Dual Haswell vs Tesla K80

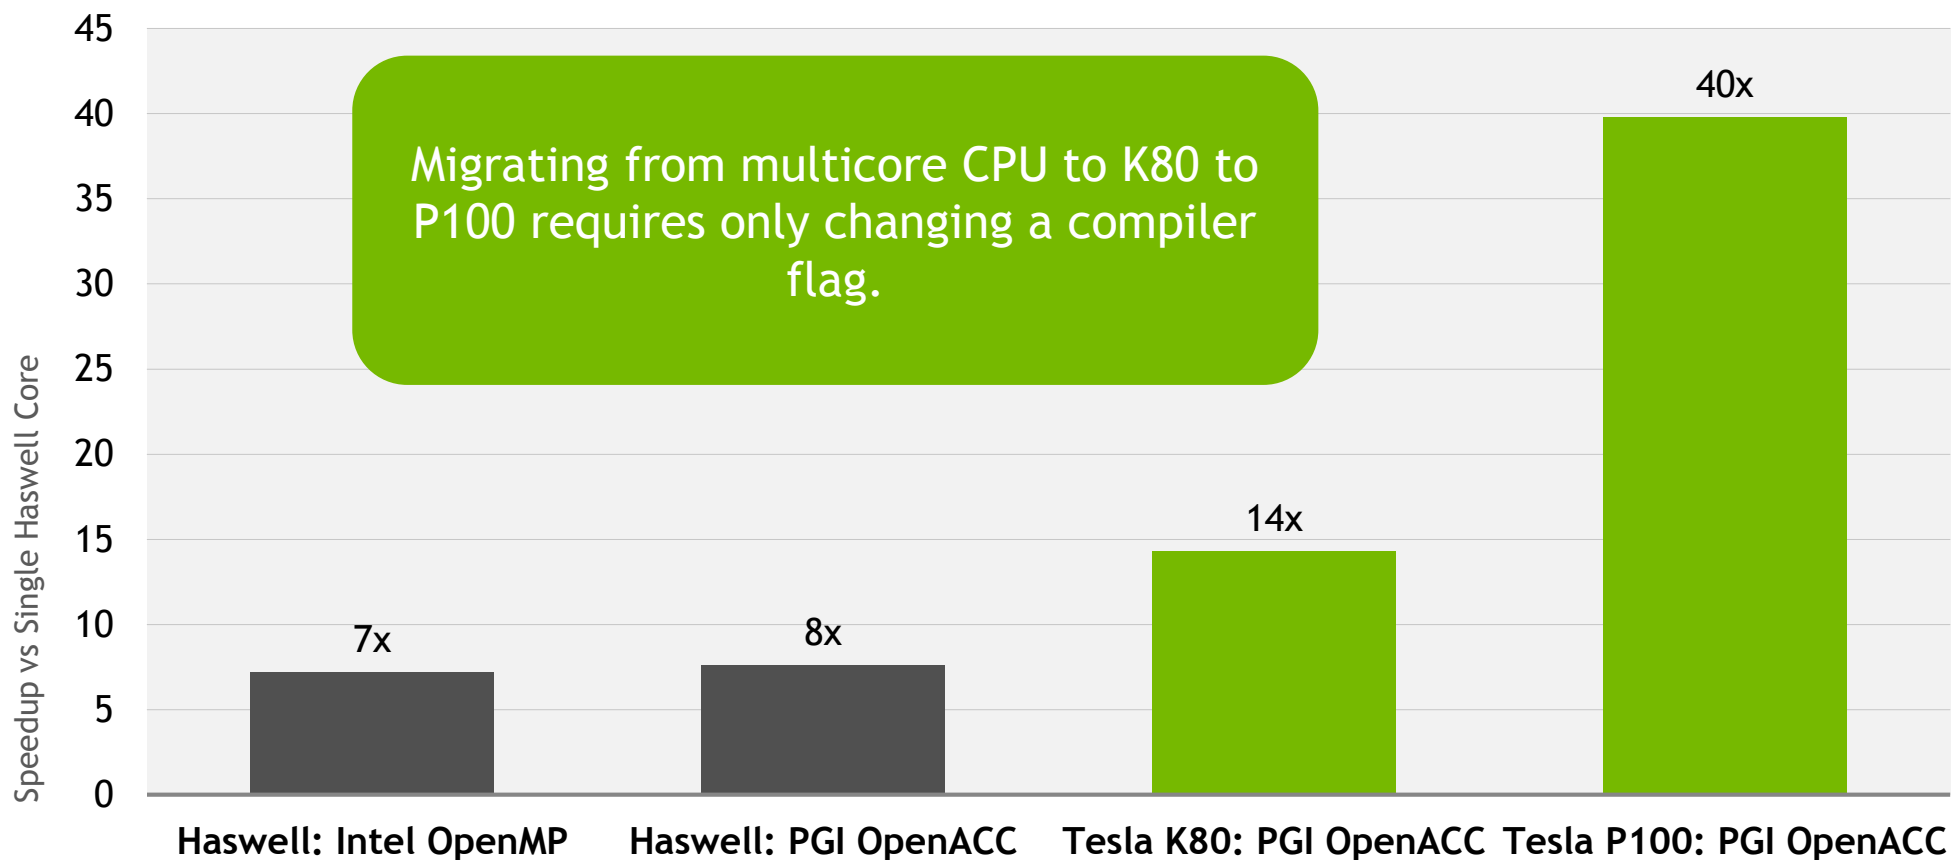


CPU: Intel Xeon E5-2698 v3, 2 sockets, 32 cores, 2.30 GHz, HT disabled

GPU: NVIDIA Tesla K80 (single GPU)

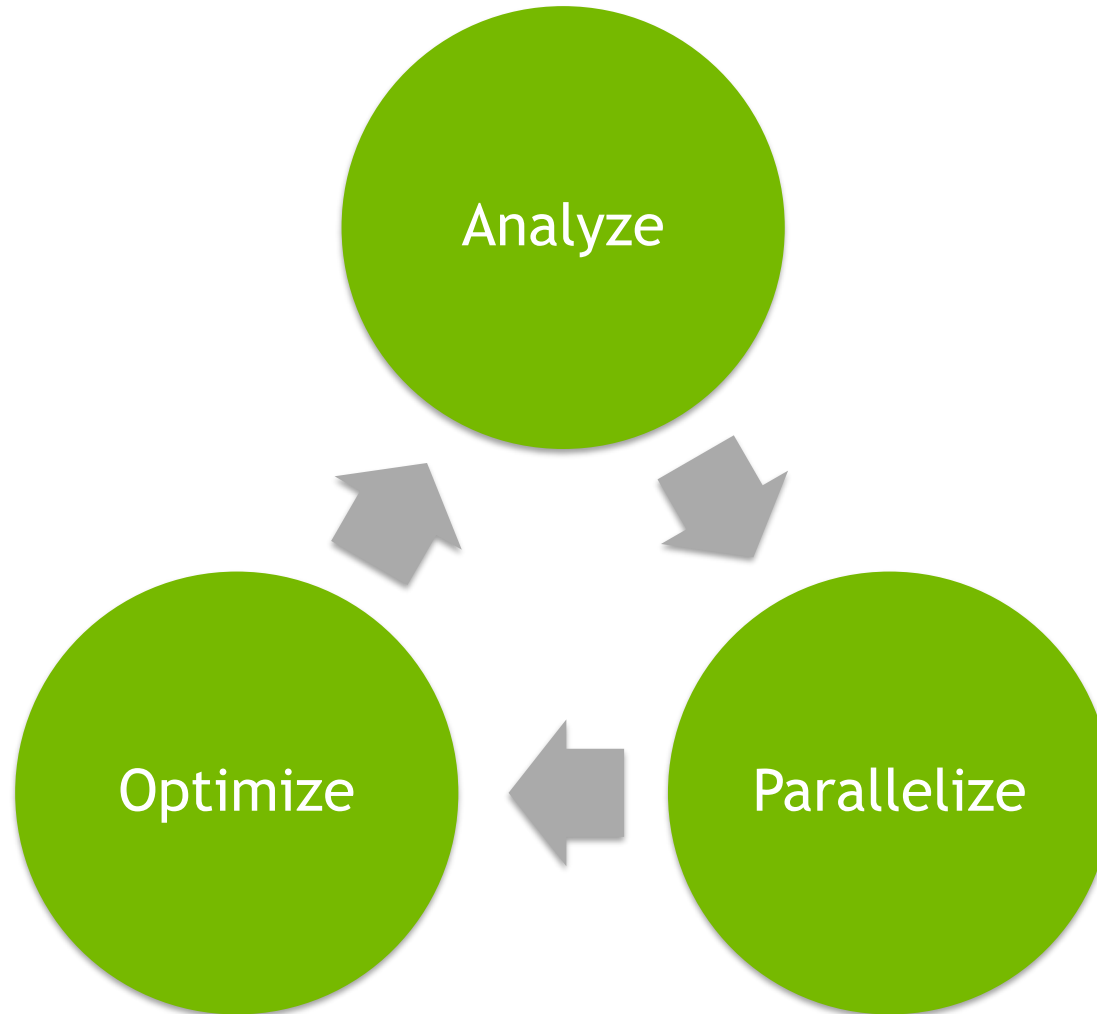
OS: CentOS 6.6, Compiler: PGI 16.5

CloverLeaf on Tesla P100 Pascal



CPU: Intel Xeon E5-2698 v3, 2 sockets, 32 cores, 2.30 GHz, HT disabled
GPU: NVIDIA Tesla K80 (single GPU), NVIDIA Tesla P100 (Single GPU)
OS: CentOS 6.6, **Compiler:** PGI 16.5

3 Steps to Accelerate with OpenACC

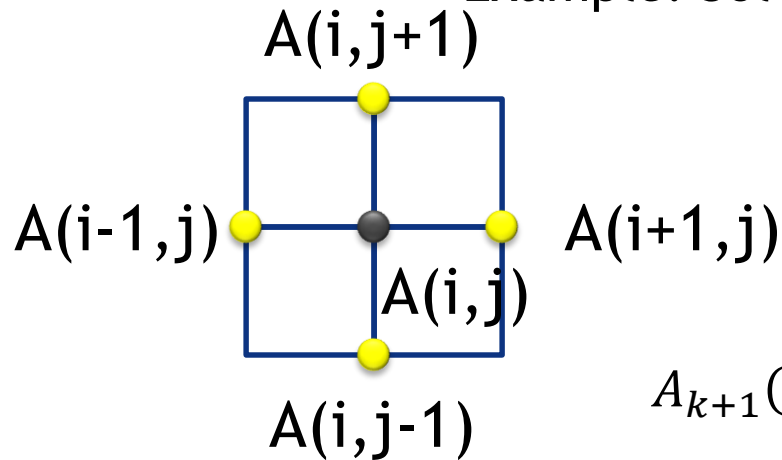


Example: Jacobi Iteration

Iteratively converges to correct value (e.g. Temperature), by computing new values at each point from the average of neighboring points.

Common, useful algorithm

Example: Solve Laplace equation in 2D: $\nabla^2 f(x, y) = 0$



$$A_{k+1}(i, j) = \frac{A_k(i-1, j) + A_k(i+1, j) + A_k(i, j-1) + A_k(i, j+1)}{4}$$

Jacobi Iteration: C Code

```
while ( err > tol && iter < iter_max ) {  
    err=0.0;  
  
    for( int j = 1; j < n-1; j++) {  
        for(int i = 1; i < m-1; i++) {  
  
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                                A[j-1][i] + A[j+1][i]);  
  
            err = max(err, abs(Anew[j][i] - A[j][i]));  
        }  
    }  
  
    for( int j = 1; j < n-1; j++) {  
        for( int i = 1; i < m-1; i++ ) {  
            A[j][i] = Anew[j][i];  
        }  
    }  
  
    iter++;  
}
```



Iterate until converged



Iterate across matrix
elements



Calculate new value from
neighbors



Compute max error for
convergence

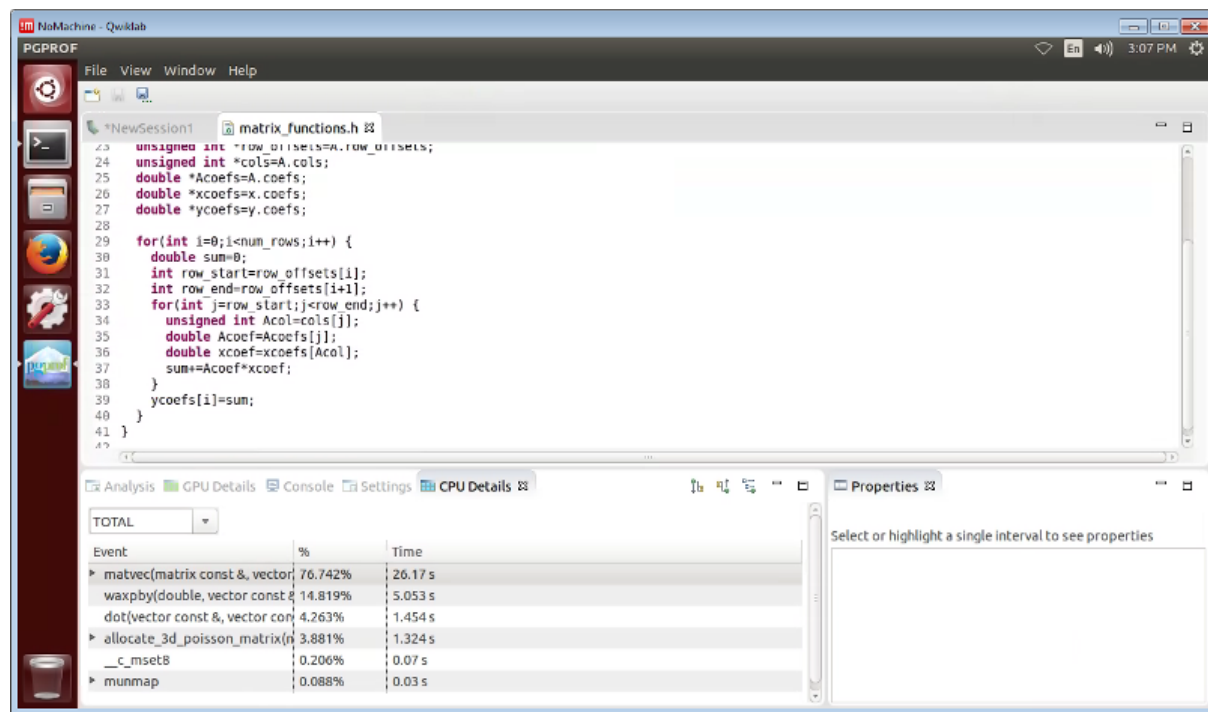


Swap input/output arrays

Analyze

Analyze

- ▶ Obtain a performance profile
- ▶ Read compiler feedback
- ▶ Understand the code



The screenshot shows the PGPROF application interface. The top pane displays a C++ code file named `matrix_functions.h` with the following content:

```
23 unsigned int *row_offsets=A.row_offsets;
24 unsigned int *cols=A.cols;
25 double *Acoefs=A.coefs;
26 double *xcoefs=x.coefs;
27 double *ycoefs=y.coefs;
28
29 for(int i=0;i<num_rows;i++) {
30     double sum=0;
31     int row_start=row_offsets[i];
32     int row_end=row_offsets[i+1];
33     for(int j=row_start;j<row_end;j++) {
34         unsigned int Acol=cols[j];
35         double Acoef=Acoefs[j];
36         double xcoef=xcoefs[Acol];
37         sum+=Acoef*xcoef;
38     }
39     ycoefs[i]=sum;
40 }
41 }
```

The bottom pane shows the CPU Details tab with a performance profile table. The table has columns for Event, %, and Time. The data is as follows:

Event	%	Time
TOTAL		
matvec(matrix const &, vector<double> const &)	76.742%	26.17 s
waxpby(double, vector const &, vector const &)	14.819%	5.053 s
dot(vector const &, vector const &)	4.263%	1.454 s
allocate_3d_poisson_matrix(n)	3.881%	1.324 s
_c_mset8	0.206%	0.07 s
munmap	0.088%	0.03 s

Obtain a Profile

A application profile helps to understand where time is spent

What routines are *hotspots*?

Focusing on the hotspots delivers the greatest performance impact

A variety of profiling tools are available: gprof, nvprof, CrayPAT, TAU, Vampir

We'll use PGProf, which comes with the PGI compiler

```
$ pgprof &
```

PGPROF Profiler

The screenshot displays the PGPROF Profiler interface. The top pane shows a C code file named `laplace2d.c` with the following code:

```
printf("Jacobi relaxation Calculation: %d x %d mesh\n", n, m);  
  
double st = omp_get_wtime();  
int iter = 0;  
  
while ( error > tol && iter < iter_max )  
{  
    error = 0.0;  
  
    for( int j = 1; j < n-1; j++)  
    {  
        Multiple markers at this line  
        - Generated 3 prefetch instructions for the loop  
        - Intensity = 1.00  
        j[i+1] + A[j][i-1]  
        + A[j-1][i] + A[j+1][i]);  
        error = fmax( error, fabs(Anew[j][i] - A[j][i]));  
    }  
  
    for( int j = 1; j < n-1; j++)  
    {  
        for( int i = 1; i < m-1; i++ )  
        {  
            A[j][i] = Anew[j][i];  
        }  
    }  
}
```

The bottom pane shows the CPU Details tab with a table of performance data:

Event	%	Time
main	67.208%	46.36 s
???	67.208%	46.36 s
__c_mcopy8	32.705%	22.56 s
__GI_memset	0.087%	0.06 s

Two green callout boxes highlight specific features:

- Inline compiler feedback.** Points to the multiple markers at the line `j[i+1] + A[j][i-1] + A[j-1][i] + A[j+1][i];`.
- Time spent by function.** Points to the table showing the time spent by different functions.

Compiler Feedback

- ▶ Before we can make changes to the code, we need to understand how the compiler is optimizing
- ▶ With PGI, this can be done with the `-Minfo` and `-Mneginfo` flags

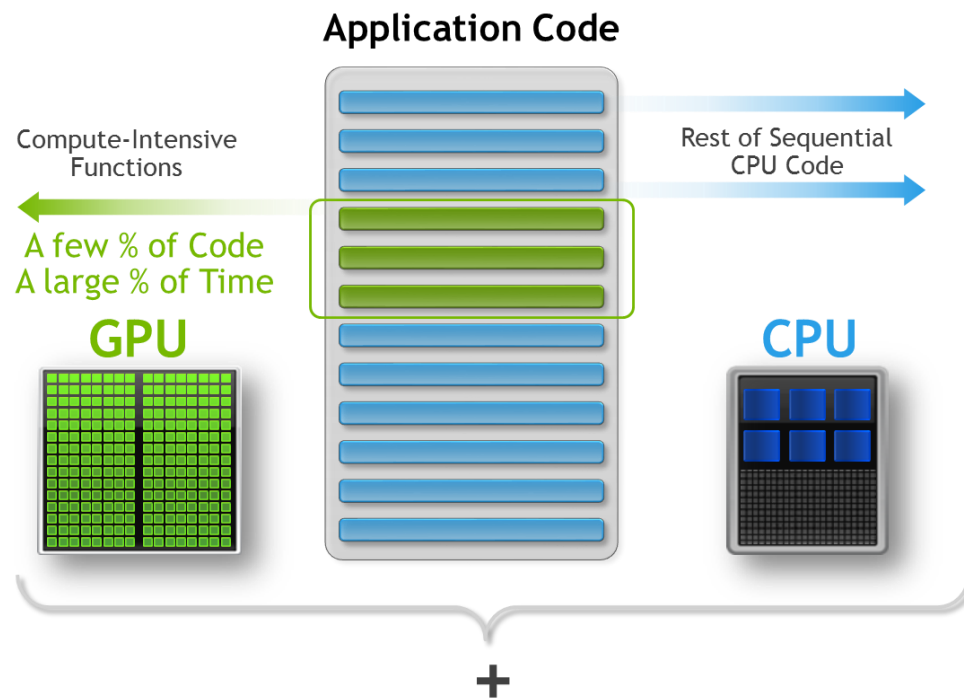
```
$ pgc++ -Minfo=all,ccff -Mneginfo
```

```
main:
    40, Loop not fused: function call before
adjacent loop
    Loop not vectorized: may not be
beneficial
    Unrolled inner loop 4 times
    Generated 3 prefetches in scalar loop
    57, Generated vector simd code for the loop
containing reductions
    Generated 3 prefetch instructions for
the loop
    67, Memory copy idiom, loop replaced by
call to __c_mcopy8
```

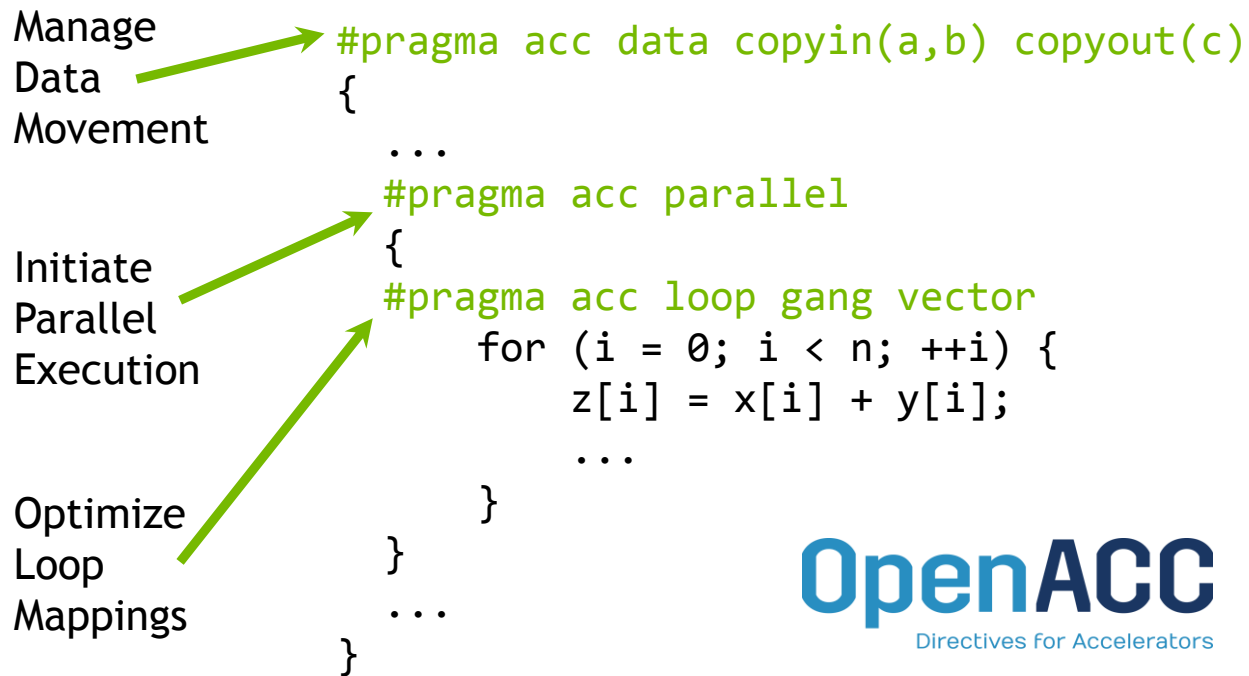

Parallelize

Parallelize

- ▶ Insert OpenACC directives around important loops
- ▶ Enable OpenACC in the compiler
- ▶ Run on a parallel platform



OpenACC Directives



- Incremental
- Single source
- Interoperable
- Performance portable
- CPU, GPU, MIC

OpenACC Parallel Directive

Generates parallelism

```
#pragma acc parallel
```

```
{
```

When encountering the *parallel* directive, the compiler will generate *1 or more parallel gangs*, which execute redundantly.

```
}
```

OpenACC Parallel Directive

Generates parallelism

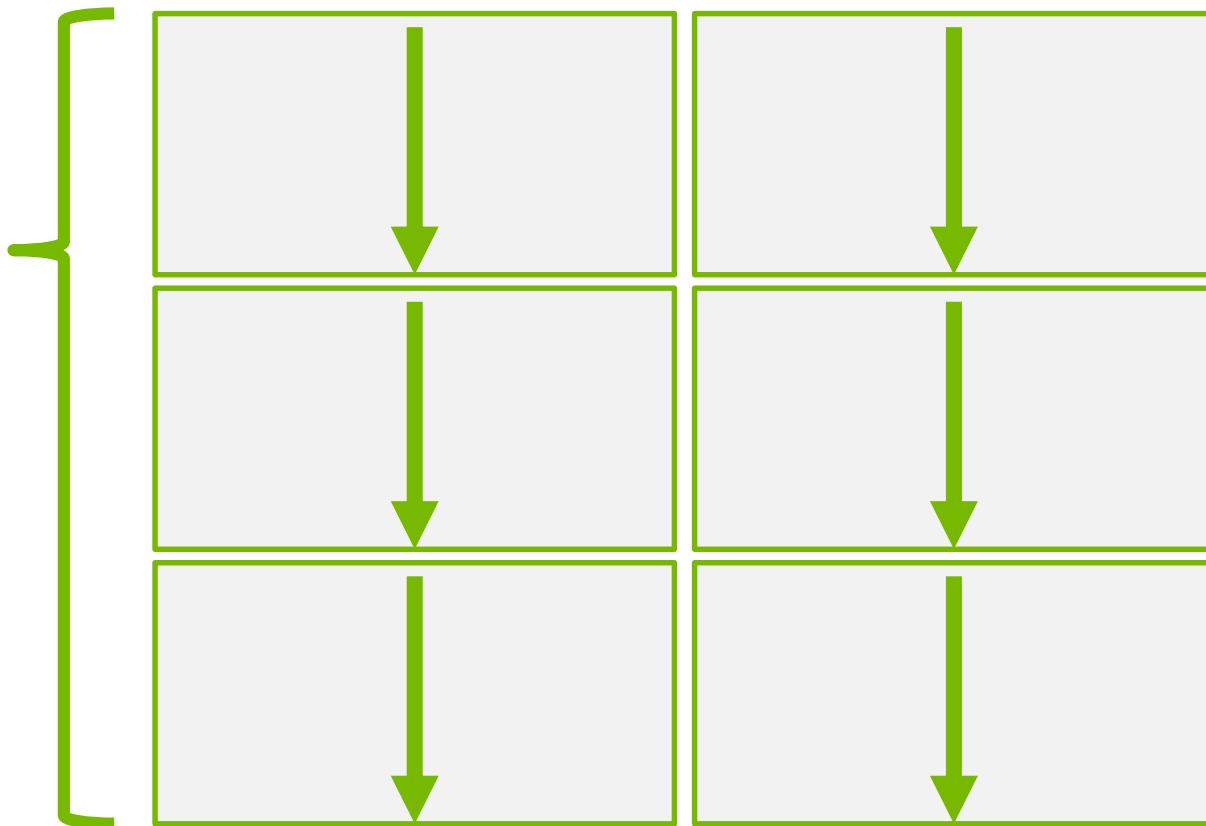
```
#pragma acc parallel
```

```
{
```



```
}
```

When encountering the *parallel* directive, the compiler will generate *1 or more parallel gangs*, which execute redundantly.



OpenACC Loop Directive

Identifies loops to run in parallel

```
#pragma acc parallel
```

```
{
```



```
}
```

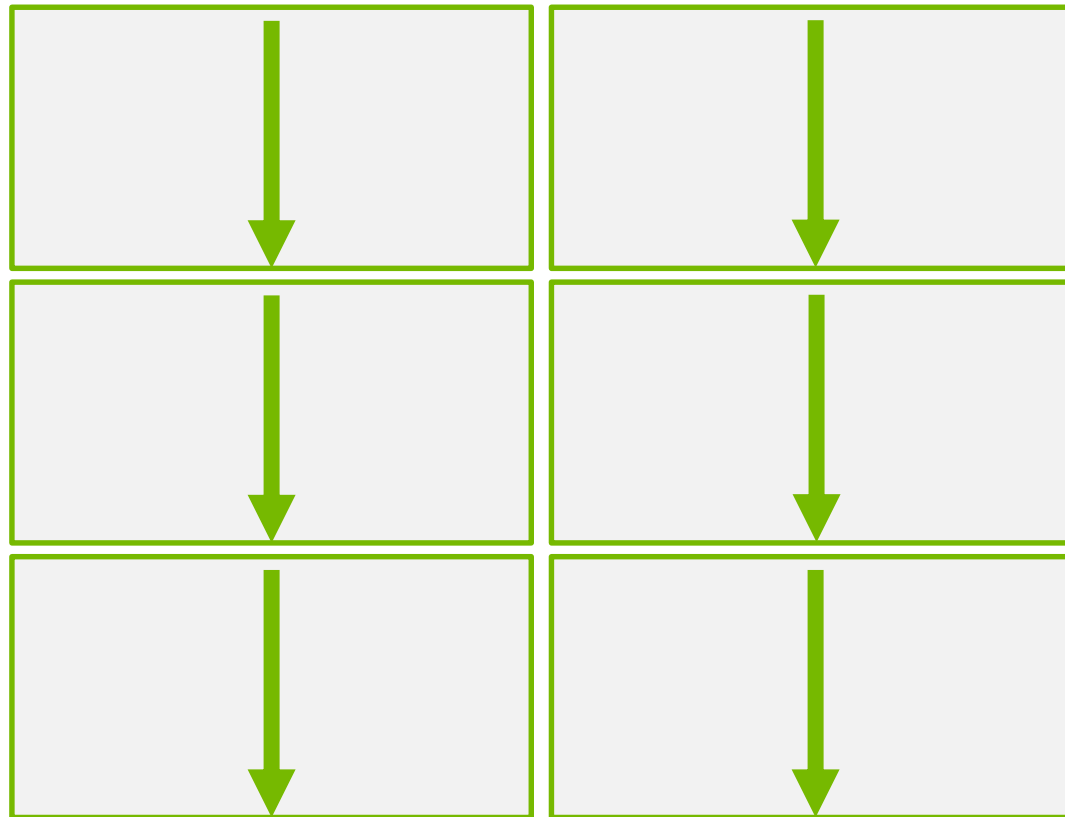
```
#pragma acc loop
```

```
for (i=0;i<N;i++)
```

```
{
```

```
}
```

The *loop* directive informs the compiler which loops to parallelize.



OpenACC Loop Directive

Identifies loops to run in parallel

```
#pragma acc parallel
```

```
{
```

```
#pragma acc loop
```

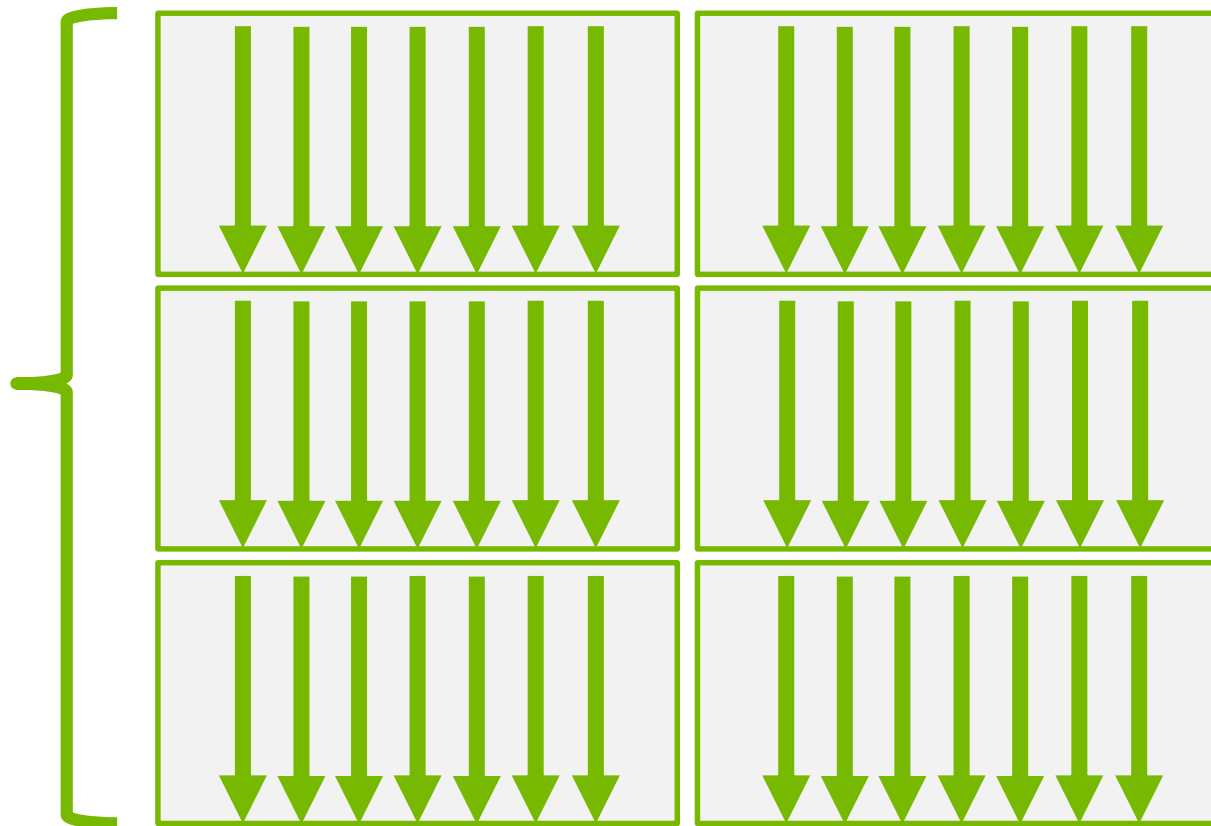
```
for (i=0;i<N;i++)
```

```
{
```

```
}
```

```
}
```

The *loop* directive
informs the compiler
which loops to
parallelize.



OpenACC Parallel Loop Directive

Generates parallelism and identifies loop in one directive

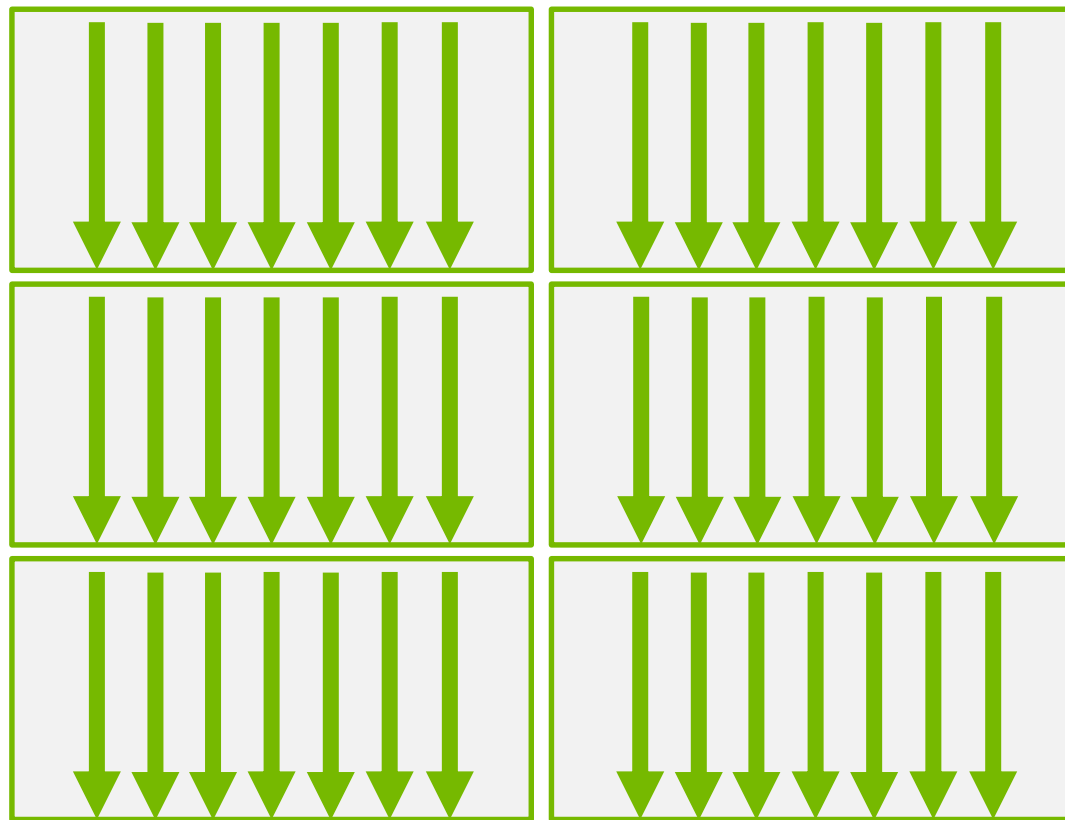
```
#pragma acc parallel loop
```

```
for (i=0;i<N;i++)
```

```
{
```

```
}
```

The *parallel* and *loop* directives are frequently combined into one.



Parallelize with OpenACC Parallel Loop

```
while ( err > tol && iter < iter_max ) {  
    err=0.0;  
  
    #pragma acc parallel loop reduction(max:err)  
    for( int j = 1; j < n-1; j++) {  
        for(int i = 1; i < m-1; i++) {  
  
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                                A[j-1][i] + A[j+1][i]);  
  
            err = max(err, abs(Anew[j][i] - A[j][i]));  
        }  
    }  
  
    #pragma acc parallel loop  
    for( int j = 1; j < n-1; j++) {  
        for( int i = 1; i < m-1; i++ ) {  
            A[j][i] = Anew[j][i];  
        }  
    }  
    iter++;  
}
```

Parallelize first loop nest,
max *reduction* required.

Parallelize second loop.

We didn't detail *how* to
parallelize the loops, just *which*
loops to parallelize.

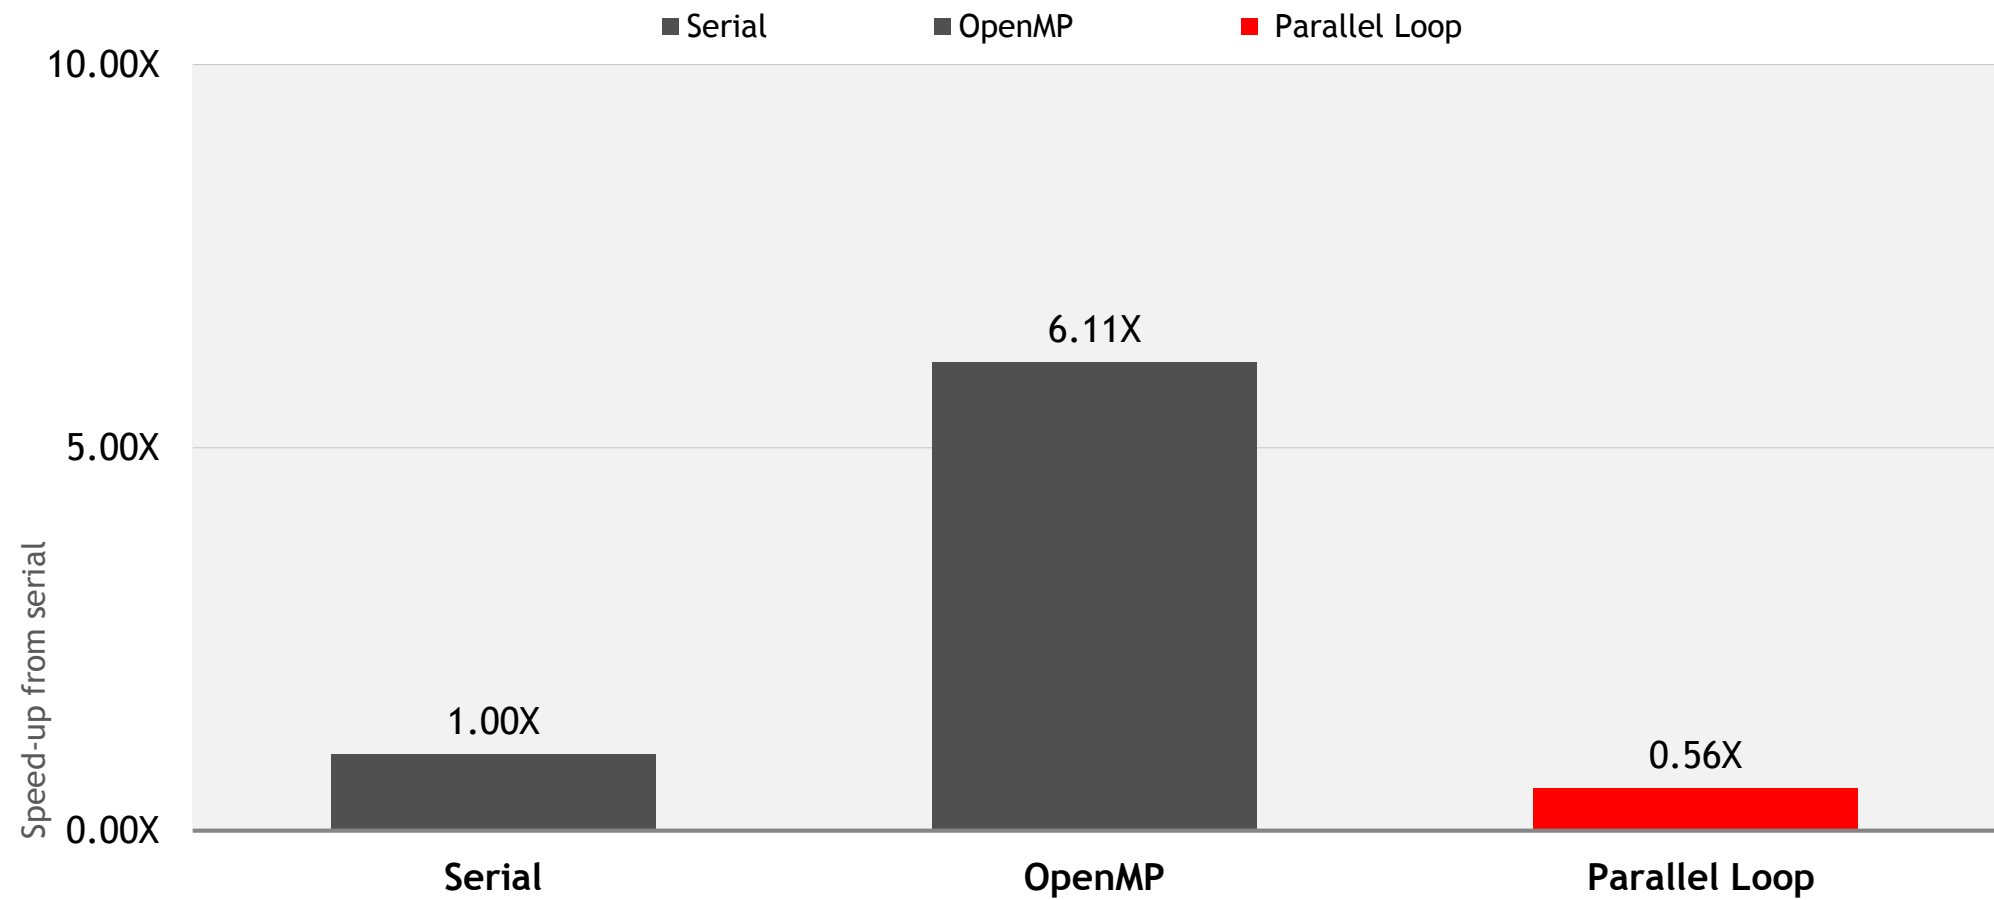
Building the code

```
$ pgcc -fast -ta=tesla,cc60 -Minfo=all,ccff laplace2d.c
```

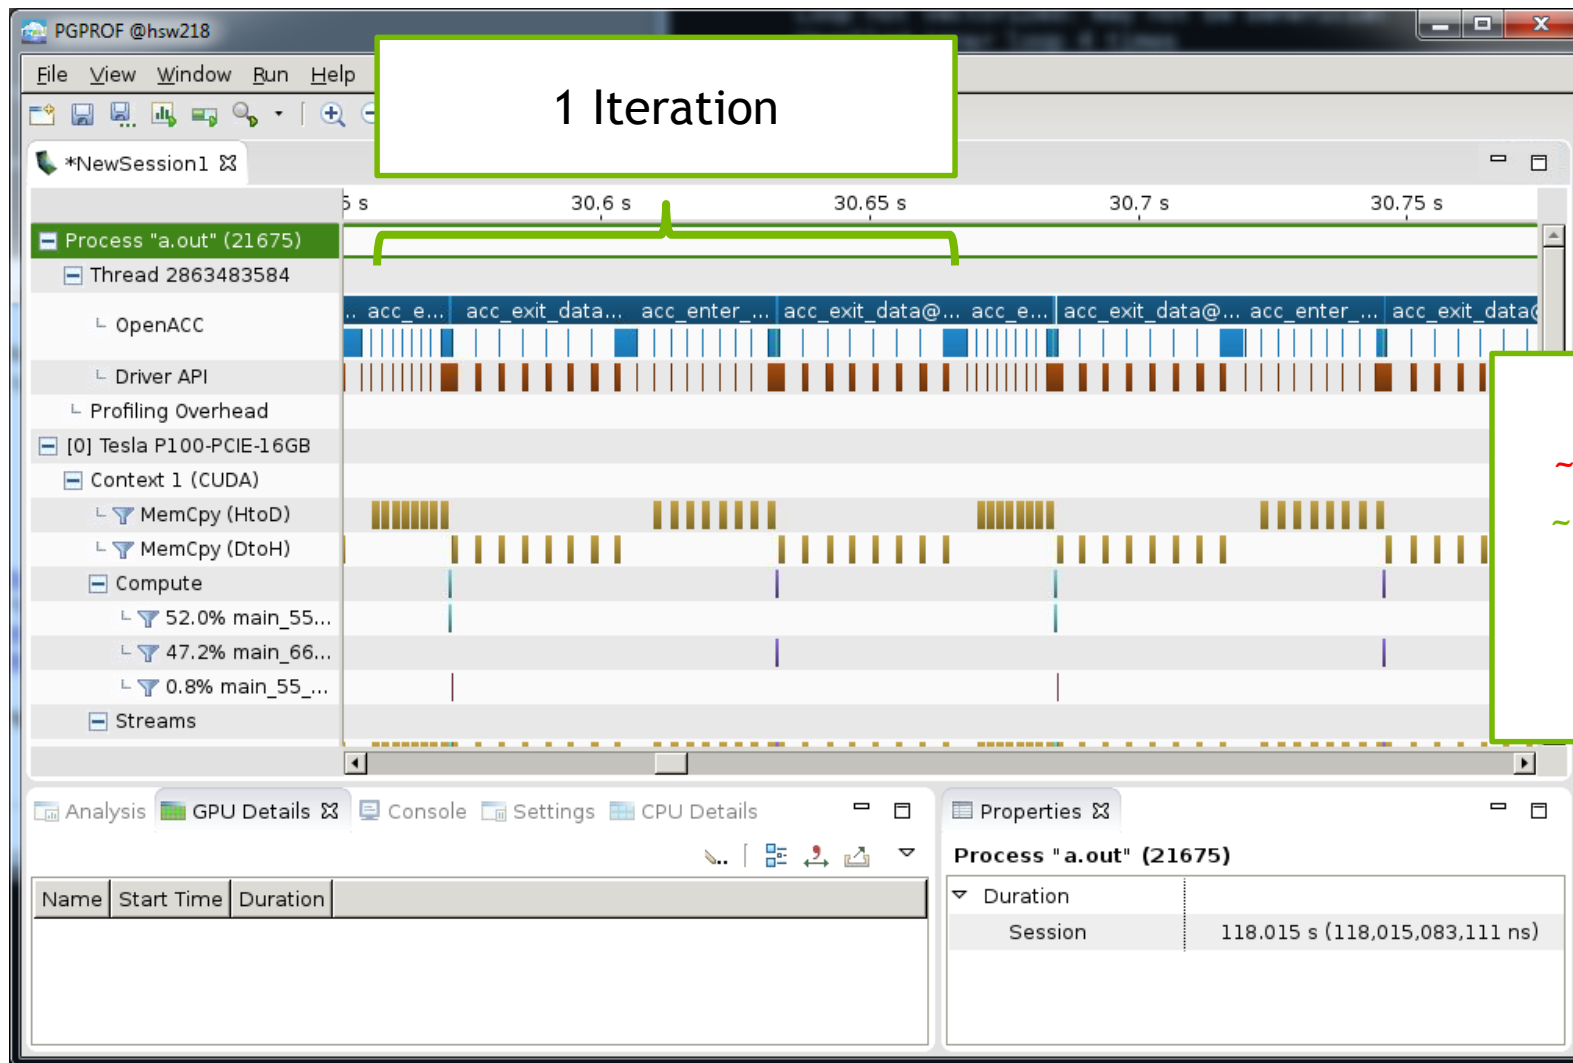
```
main:
```

```
40, Loop not fused: function call before adjacent loop
    Loop not vectorized: may not be beneficial
    Unrolled inner loop 4 times
    Generated 3 prefetches in scalar loop
51, Loop not vectorized/parallelized: potential early exits
55, Accelerator kernel generated
    Generating Tesla code
    55, Generating reduction(max:error)
    56, #pragma acc loop gang /* blockIdx.x */
    58, #pragma acc loop vector(128) /* threadIdx.x */
55, Generating implicit copyout(Anew[1:4094][1:4094])
    Generating implicit copyin(A[:][:])
58, Loop is parallelizable
66, Accelerator kernel generated
    Generating Tesla code
    67, #pragma acc loop gang /* blockIdx.x */
    69, #pragma acc loop vector(128) /* threadIdx.x */
66, Generating implicit copyin(Anew[1:4094][1:4094])
    Generating implicit copyout(A[1:4094][1:4094])
69, Loop is parallelizable
```

OpenACC Performance (Step 1)



What went wrong?



1 Iteration

Per-iteration:

~450ms Copying Data

~1.1ms Computing on GPU

Why?

Step 1 Compiler Feedback

```
51 while ( error > tol && iter < iter_max )
52 {
53     error = 0.0;
54
55     #pragma acc parallel loop
56         reduction(max:error)
57     for( int j=1; j < n-1; j++)
58     {
59         for(int i=1; i < m-1; i++)
60         {
61             ...
62         }
63     }
64 }
65-76 ...
77 iter++;
78 }
```

```
main:
51, Loop not vectorized/parallelized:
potential early exits
55, Accelerator kernel generated
Generating Tesla code
55, Generating reduction(max:error)
56, #pragma acc loop gang /*
blockIdx.x */
58, #pragma acc loop vector(128) /*
threadIdx.x */
55, Generating implicit copyin(A[:][:])
Generating implicit
copyout(Anew[1:4094][1:4094])
```

The compiler implicitly copies A and Anew to/from the GPU in case we need them, but do we?

Optimizing Data Movement

```
while ( err > tol && iter < iter_max ) {  
    err=0.0;  
  
    #pragma acc parallel loop  
    for( int j = 1; j < n-1; j++) {  
        for(int i = 1; i < m-1; i++) {  
  
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                                A[j-1][i] + A[j+1][i]);  
  
            err = max(err, abs(Anew[j][i] - A[j][i]));  
        }  
    }  
  
    #pragma acc parallel loop  
    for( int j = 1; j < n-1; j++) {  
        for( int i = 1; i < m-1; i++ ) {  
            A[j][i] = Anew[j][i];  
        }  
    }  
  
    iter++;  
}
```

Does the CPU need the data
between these loop nests?

Does the CPU need the data
between iterations of the
convergence loop?

Optimize

Optimize

- ✓ ▶ Get new performance data from parallel execution
- ▶ Remove unnecessary data transfer to/from GPU
- ▶ Guide the compiler to better loop decomposition



Structured Data Regions

The **data** directive defines a region of code in which GPU arrays remain on the GPU and are shared among all kernels in that region.

```
#pragma acc data
{
#pragma acc parallel loop
...

#pragma acc parallel loop
...
}
```

Data Region

Arrays used within the data region will remain on the GPU until the end of the data region.

Data Clauses

`copyin (list)`

Allocates memory on GPU and copies data from host to GPU when entering region.

`copyout (list)`

Allocates memory on GPU and copies data to the host when exiting region.

`copy (list)`

Allocates memory on GPU and copies data from host to GPU when entering region and copies data to the host when exiting region. (Structured Only)

`create (list)`

Allocates memory on GPU but does not copy.

`delete(list)`

Deallocate memory on the GPU without copying. (Unstructured Only)

`present (list)`

Data is already present on GPU from another containing data region.

(!) All of these will check if the data is already present first and reuse if found.

Optimized Data Movement


```
#pragma acc data copy(A) create(Anew)
while ( err > tol && iter < iter_max ) {
    err=0.0;

    #pragma acc parallel loop reduction(max:err)
    for( int j = 1; j < n-1; j++) {
        for(int i = 1; i < m-1; i++) {

            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +
                                A[j-1][i] + A[j+1][i]);

            err = max(err, abs(Anew[j][i] - A[j][i]));
        }
    }

    #pragma acc parallel loop
    for( int j = 1; j < n-1; j++) {
        for( int i = 1; i < m-1; i++ ) {
            A[j][i] = Anew[j][i];
        }
    }
    iter++;
}
```




Copy A to/from the
accelerator only when
needed.

Create Anew as a device
temporary.

Rebuild the Code

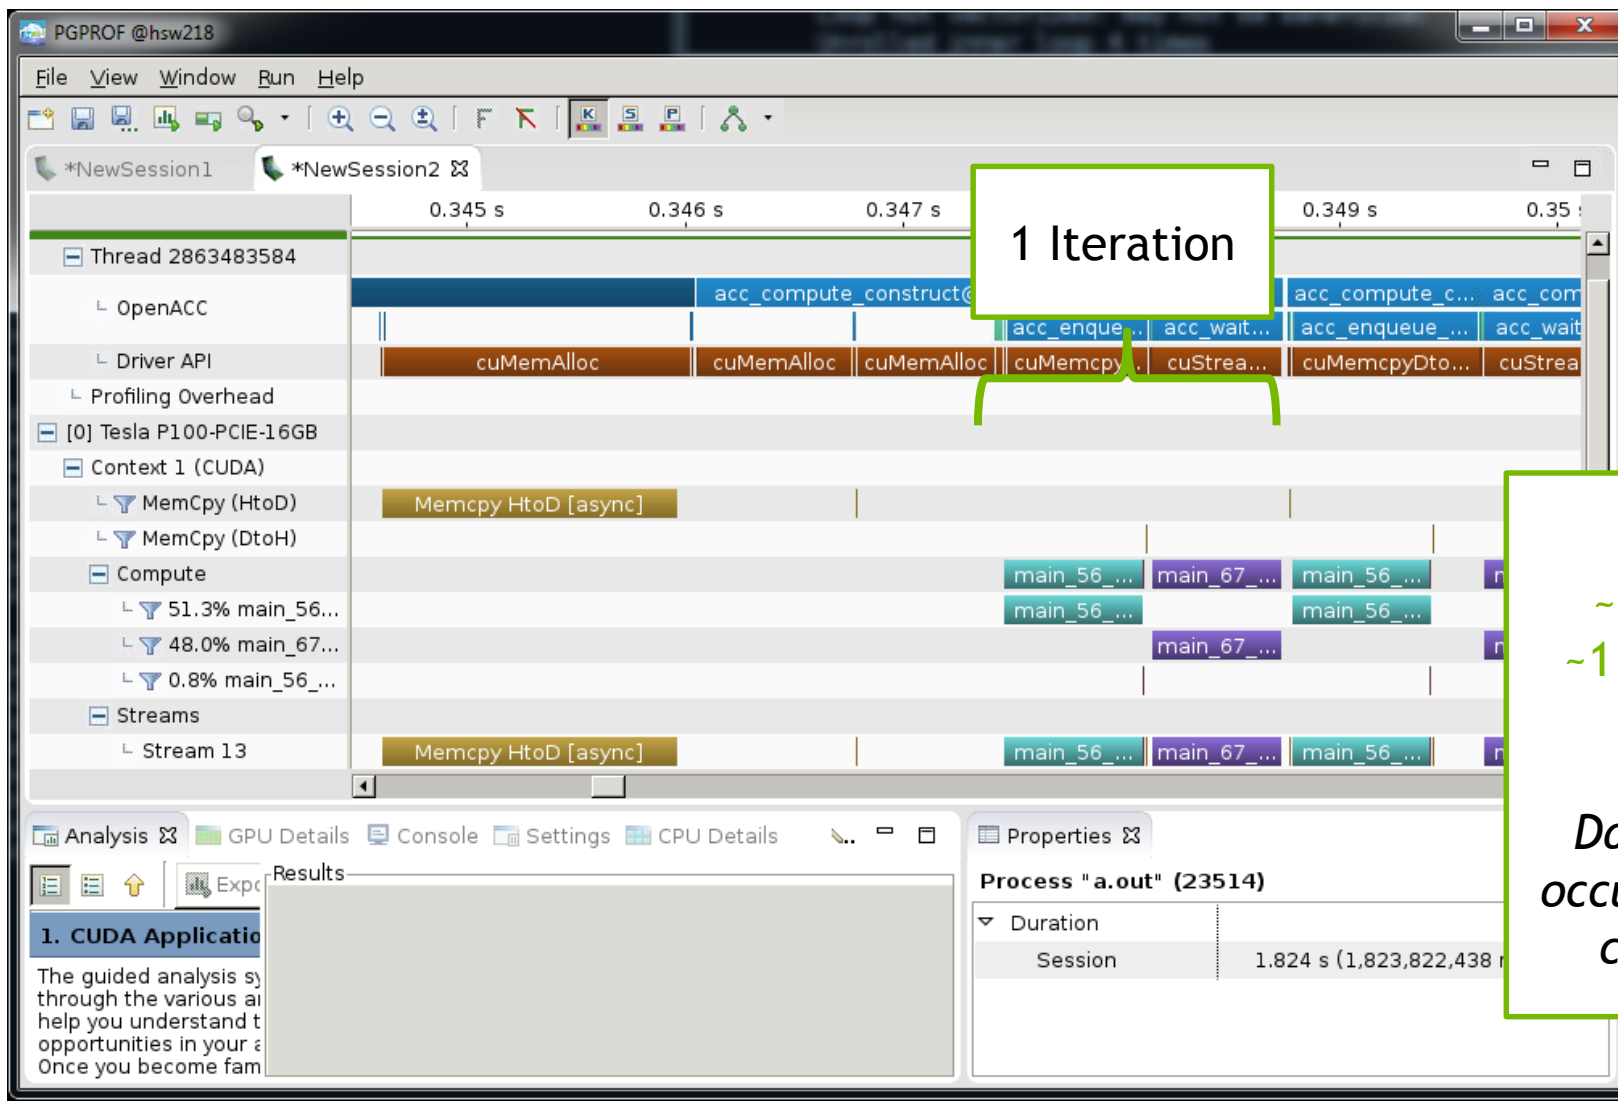
main:

```
51, Generating create(Anew[:][:])  
    Generating copy(A[:][:])  
56, Accelerator kernel generated  
    Generating Tesla code  
    56, Generating reduction(max:error)  
    57, #pragma acc loop gang /* blockIdx.x */  
    59, #pragma acc loop vector(128) /* threadIdx.x */  
59, Loop is parallelizable  
67, Accelerator kernel generated  
    Generating Tesla code  
    68, #pragma acc loop gang /* blockIdx.x */  
    70, #pragma acc loop vector(128) /* threadIdx.x */  
70, Loop is parallelizable
```



Now data movement only
happens at our data
region.

PGProf Timeline Now

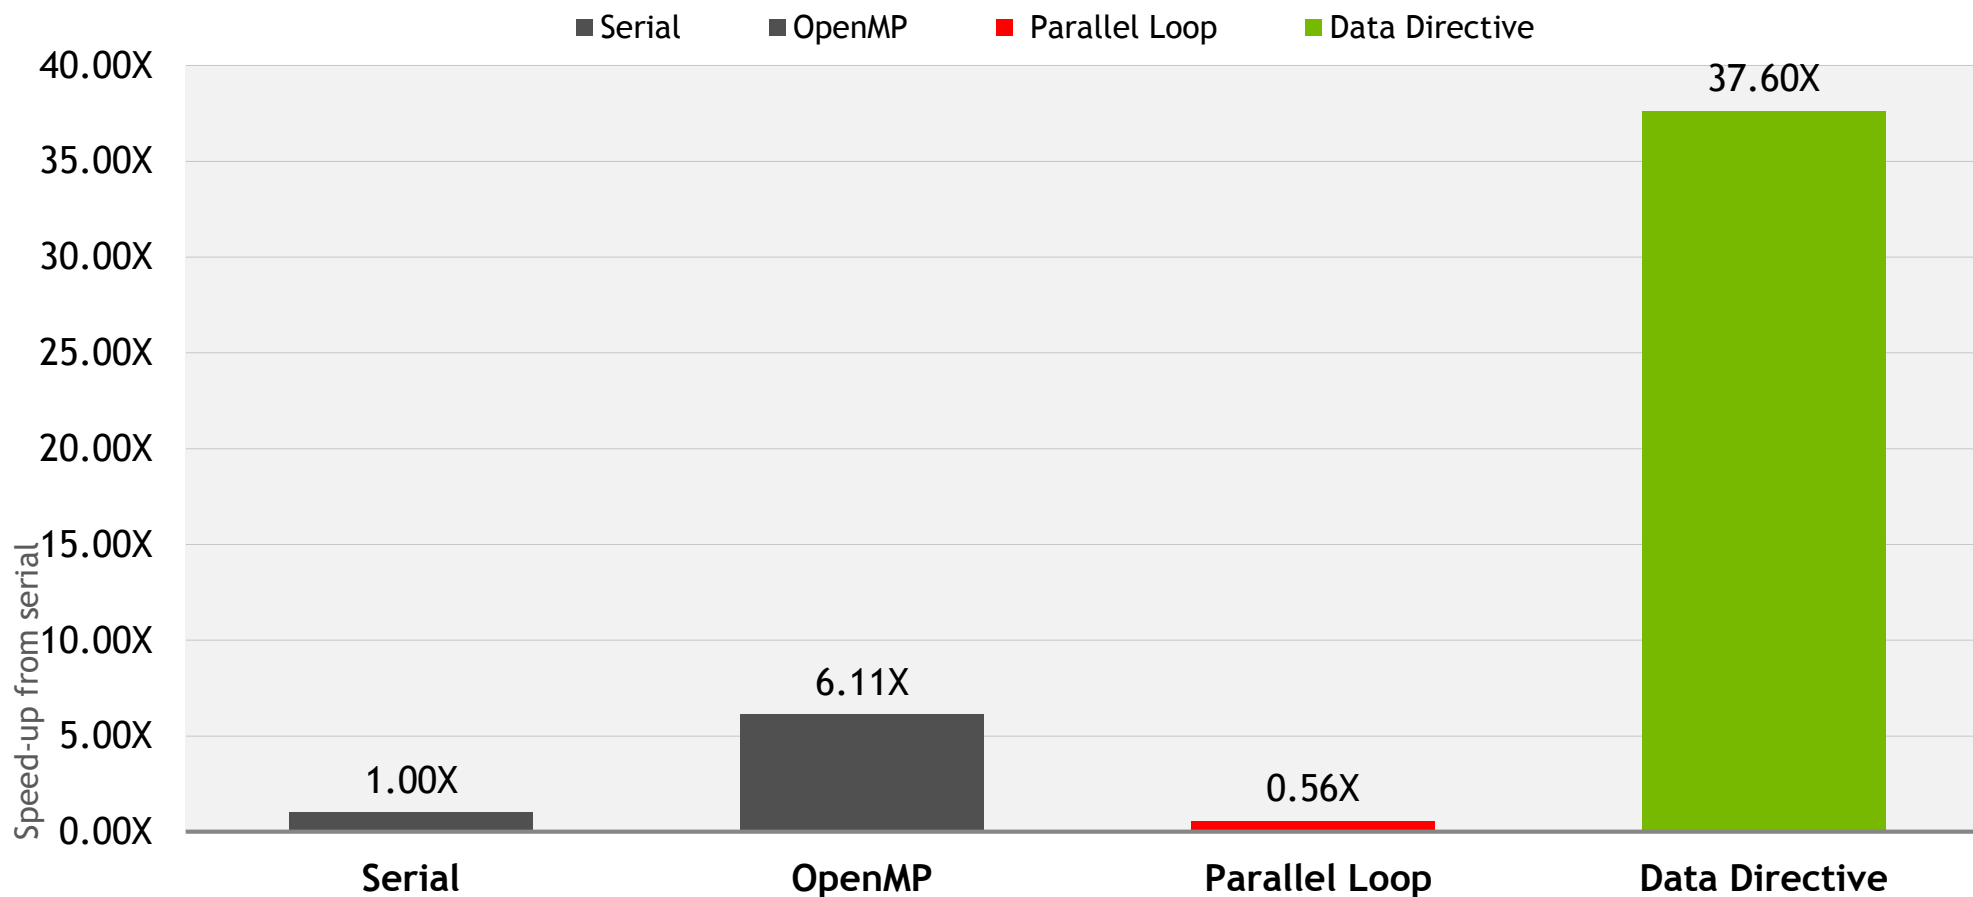


1 Iteration

Per-iteration:
~0ms Copying Data
~1.1ms Computing on GPU

Data movement now occurs before and after convergence loop.

OpenACC Performance (Step 2)



Array Shaping

Compiler sometimes cannot determine size of arrays

Must specify explicitly using data clauses and array “shape”

Partial arrays must be contiguous

C/C++

```
#pragma acc data copyin(a[0:N]) copyout(b[start:count])
```

Fortran

```
!$acc data copyin(a(1:N)) copyout(b(start:start+count))
```

Unstructured Data: C++ Classes

- ▶ Unstructured Data Regions enable OpenACC to be used in C++ classes
- ▶ Unstructured data regions can be used whenever data is allocated and initialized in a different scope than where it is freed (e.g. Fortran modules).

```
class Matrix {  
    Matrix(int n) {  
        len = n;  
        v = new double[len];  
        #pragma acc enter data  
            create(v[0:len])  
    }  
    ~Matrix() {  
        #pragma acc exit data  
            delete(v[0:len])  
        delete[] v;  
    }  
  
private:  
    double* v;  
    int len;  
};
```


OpenACC Update Directive

Programmer specifies an array (or part of an array) that should be refreshed within a data region.

```
#pragma acc data create(a)
{
do_something_on_device()
#pragma acc update self(a)
do_something_on_host()
#pragma acc update device(a)
}
```



Copy “a” from GPU to
CPU



Copy “a” from CPU to
GPU

Optimize Loops

Now let's look at how our iterations get mapped to hardware.

Compilers give their best guess about how to transform loops into parallel kernels, but sometimes they need more information.

This information could be our knowledge of the code or based on profiling.

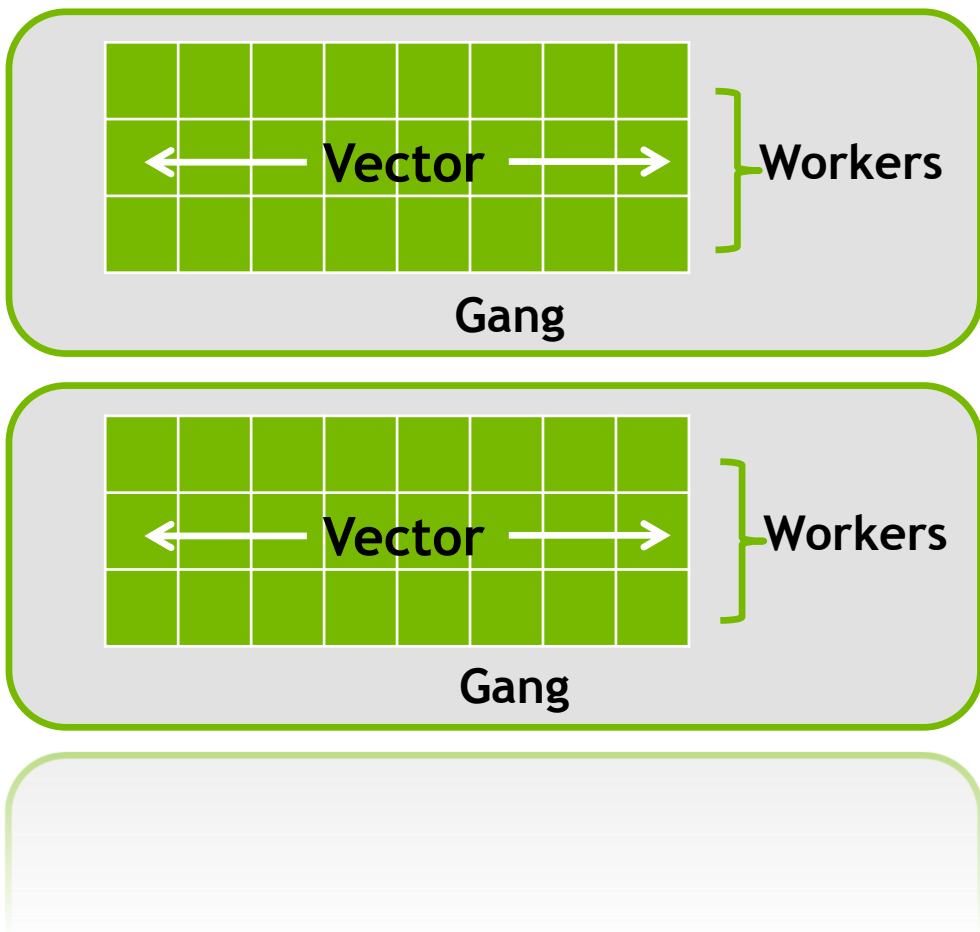
Step 2 Compiler Feedback

```
51 #pragma acc data copy(A) create(Anew)
52 while ( error > tol && iter < iter_max )
53 {
54     error = 0.0;
55
56 #pragma acc parallel loop
57     reduction(max:error)
58     for( int j=1; j < n-1; j++)
59     {
60         for(int i=1; i < m-1; i++)
61         {
62             ...
63         }
64     }
65 }
66-77 ...
78 iter
79 }
```

The compiler is *vectorizing* the inner loops and breaking the outer loops across *gangs*.

```
main:
51, Generating create(Anew[:][:])
    Generating copy(A[:][:])
56, Accelerator kernel generated
    Generating Tesla code
56, Generating reduction(max:error)
57, #pragma acc loop gang /*
    blockIdx.x */
59, #pragma acc loop vector(128) /*
    threadIdx.x */
59, Loop is parallelizable
67, Accelerator kernel generated
    Generating Tesla code
68, #pragma acc loop gang /*
    blockIdx.x */
70, #pragma acc loop vector(128) /*
    threadIdx.x */
70, Loop is parallelizable
```

OpenACC: 3 Levels of Parallelism



- *Vector* threads work in lockstep (SIMD/SIMT parallelism)
- *Workers* compute a vector
- *Gangs* have 1 or more workers and share resources (such as cache, the streaming multiprocessor, etc.)
- Multiple gangs work independently of each other

The loop Directive

The **loop** directive gives the compiler additional information about the *next* loop in the source code through several clauses.

- **independent** - all iterations of the loop are independent
- **auto** - instructs the compiler to analyze the loop
- **collapse (N)** - turn the next N loops into one, flattened loop
- **tile (N[,M,...])** - break the next 1 or more loops into *tiles* based on the provided dimensions.
- **gang, worker, vector, seq** - Describes how to parallelize the loop.

OpenACC gang, worker, vector Clauses

gang, **worker**, and **vector** can be added to a loop clause

A parallel region can only specify one of each gang, worker, vector

Control the size using the following clauses on the parallel region

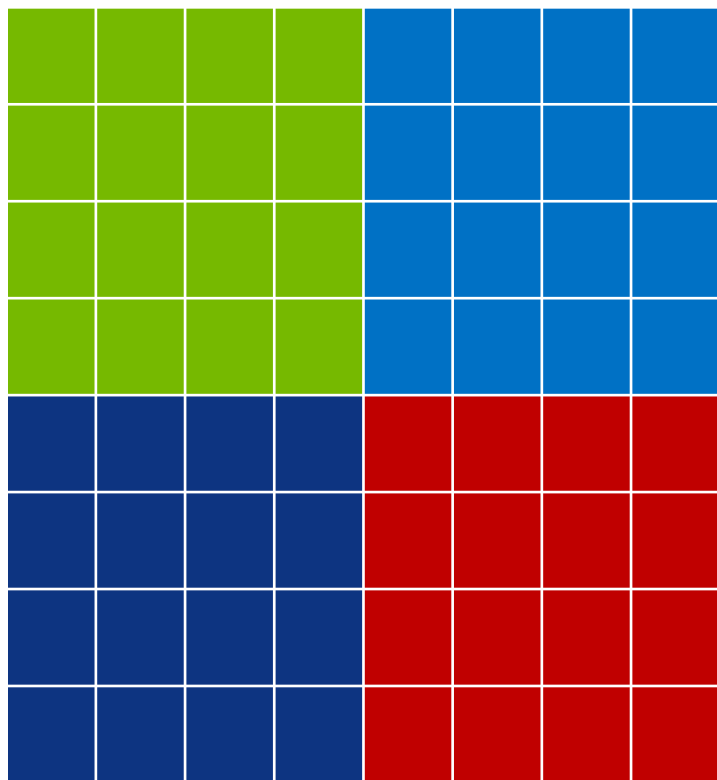
num_gangs(n), **num_workers(n)**, **vector_length(n)**

```
#pragma acc parallel loop gang
for (int i = 0; i < n; ++i)
    #pragma acc loop vector
    for (int j = 0; j < n; ++j)
        ...
```

```
#pragma acc parallel vector_length(32)
#pragma acc loop gang worker
for (int i = 0; i < n; ++i)
    #pragma acc loop vector
    for (int j = 0; j < n; ++j)
        ...
```

The tile Clause

Operate on smaller blocks of the operation to exploit data locality



```
#pragma acc loop tile(4,4)
for(i = 1; i <= ROWS; i++) {
    for(j = 1; j <= COLUMNS; j++) {
        Temp[i][j] = 0.25 *
            (Temp_last[i+1][j] +
             Temp_last[i-1][j] +
             Temp_last[i][j+1] +
             Temp_last[i][j-1]);
    }
}
```

Tile to Exploit Locality

```
#pragma acc data copy(A) create(Anew)
while ( err > tol && iter < iter_max ) {
    err=0.0;
    #pragma acc parallel loop reduction(max:err) tile(32,16)
    for( int j = 1; j < n-1; j++) {
        for(int i = 1; i < m-1; i++) {

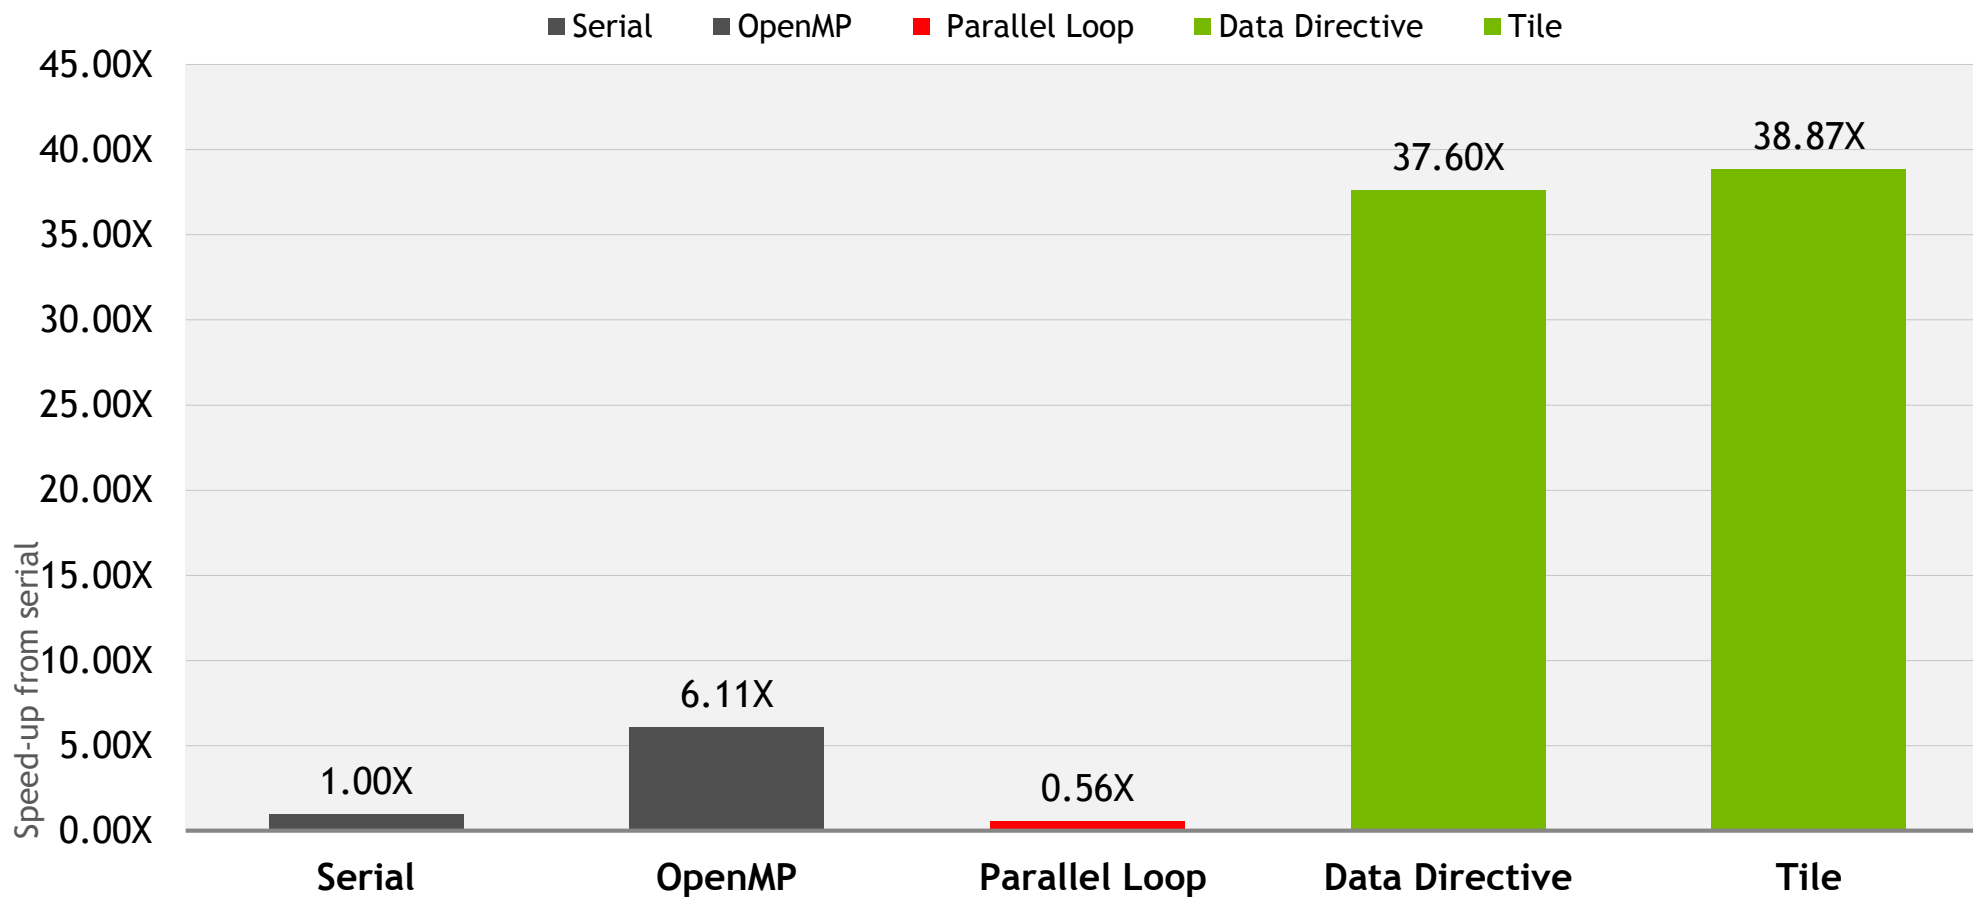
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +
                                A[j-1][i] + A[j+1][i]);

            err = max(err, abs(Anew[j][i] - A[j][i]));
        }
    }

    #pragma acc parallel loop tile(32,16)
    for( int j = 1; j < n-1; j++) {
        for( int i = 1; i < m-1; i++ ) {
            A[j][i] = Anew[j][i];
        }
    }
    iter++;
}
```

Through experimentation I found that 32x16 tiles gave me the best performance.

OpenACC Performance (Final)



OpenACC kernels Directive

Identifies a region of code where I think the compiler can turn *loops* into *kernels*

```
#pragma acc kernels
```

```
{  
for(int i=0; i<N; i++)  
{  
    x[i] = 1.0;  
    y[i] = 2.0;  
}
```

} kernel 1

```
for(int i=0; i<N; i++)  
{  
    y[i] = a*x[i] + y[i];  
}  
}
```

} kernel 2

The compiler identifies
2 parallel loops and
generates 2 kernels.

OpenACC parallel loop vs. kernels

PARALLEL LOOP

- ▶ Programmer's responsibility to ensure safe parallelism
- ▶ Will parallelize what a compiler may miss
- ▶ Straightforward path from OpenMP

KERNELS

- ▶ Compiler's responsibility to analyze the code and parallelize what is safe.
- ▶ Can cover larger area of code with single directive
- ▶ Gives compiler additional leeway to optimize.
- ▶ Compiler sometimes gets it wrong.

Both approaches are equally valid and can perform equally well.

In Closing

Where to find help

- OpenACC Course Recordings - <https://developer.nvidia.com/openacc-courses>
- PGI Website - <http://www.pgroup.com/resources>
- OpenACC on StackOverflow - <http://stackoverflow.com/questions/tagged/openacc>
- PGI Community Edition - <http://www.pgroup.com/products/community.htm>
- Parallel Forall Blog - <http://devblogs.nvidia.com/parallelforall/>
- GPU Technology Conference - <http://www.gputechconf.com/>
- OpenACC Website - <http://openacc.org/>

Questions? Email openacc@nvidia.com



QWIKLABS - Hands on in the cloud

QWIKLAB

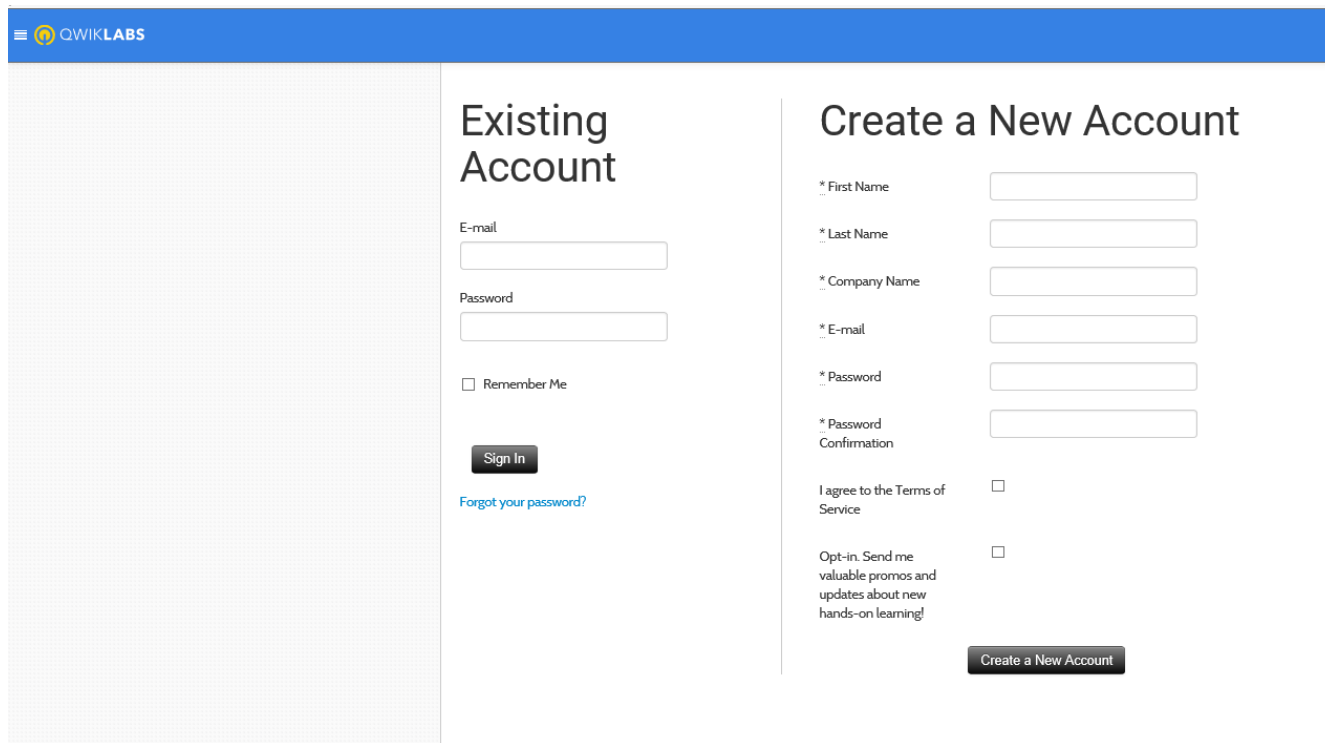
<http://nvlabs.qwiklab.com>

Navigate and login to qwiklab

Start lab:

OpenACC - 2x in 4 Steps

Complete lab



The screenshot shows the QwikLab website interface. At the top is a blue header with the QwikLab logo. The main content area is divided into two columns. The left column, titled 'Existing Account', contains a login form with fields for 'E-mail' and 'Password', a 'Remember Me' checkbox, a 'Sign In' button, and a link for 'Forgot your password?'. The right column, titled 'Create a New Account', contains a registration form with fields for 'First Name', 'Last Name', 'Company Name', 'E-mail', 'Password', and 'Password Confirmation'. Below these fields are two checkboxes: 'I agree to the Terms of Service' and 'Opt-in. Send me valuable promos and updates about new hands-on learning!'. At the bottom of the right column is a 'Create a New Account' button.

QWIKLABS

Existing Account

E-mail

Password

☐ Remember Me

[Sign In](#)

[Forgot your password?](#)

Create a New Account

* First Name

* Last Name

* Company Name

* E-mail

* Password

* Password Confirmation

I agree to the Terms of Service ☐

Opt-in. Send me valuable promos and updates about new hands-on learning! ☐

[Create a New Account](#)